

Games and the Minimax algorithm

April 1, 2021

Games and the Minimax algorithm

- ▶ We will discuss some aspects of Game Theory

Games and the Minimax algorithm

- ▶ We will discuss some aspects of Game Theory
- ▶ In particular we will discuss the Minimax Algorithm

Games and the Minimax algorithm

- ▶ We will discuss some aspects of Game Theory
- ▶ In particular we will discuss the Minimax Algorithm
- ▶ The concepts will be applied to simple examples.

Overview

Introduction to Game Theory

- Normal form description

- Optimal strategies

The Minimax Algorithm

- Sequential games

- Concept of recursion

- Minimax algorithm

- Limitations of the Minimax algorithm

- Alpha-Beta pruning

- Intermediate scoring

Formalization of games

- ▶ Games can be formalized mathematically

Zero sum normal form games

- ▶ We will consider games with two players.

Zero sum normal form games

- ▶ We will consider games with two players.
- ▶ Both players play simultaneously. Player A does action a , players B does action b .

Zero sum normal form games

- ▶ We will consider games with two players.
- ▶ Both players play simultaneously. Player A does action a , players B does action b .
- ▶ Their action results in a gain $g(a, b)$.

Zero sum normal form games

- ▶ We will consider games with two players.
- ▶ Both players play simultaneously. Player A does action a , player B does action b .
- ▶ Their action results in a gain $g(a, b)$.
- ▶ We assume that what A wins is what B loses : hence the term **zero-sum** game.
- ▶ For instance A wins $g(a, b)$ and B "wins" $-g(a, b)$.

Examples

- ▶ Paper, Scissors, Stone
- ▶ football penalty

Normal form description

- ▶ These games can be represented by a **payoff matrix**.

Zero sum normal form games

		Player B action	
		a	b
Player A action	a	$g(a, b)$	$g(a, b)$
	b	$g(a, b)$	$g(a, b)$

Table: In this game, the players can perform two possible actions. Since the game is zero-sum, it is sufficient to represent the gain of player A .

Example

		Player B action	
		a	b
Player A action	a	2	3
	b	1	8

Table: Example gains.

Concept of *supinf*

		Player B		
		a	b	c
Player A	a	2	0	9
	b	4	4	7
	c	10	1	3

Table: What is the gain the A can be **sure** of obtaining ?

Concept of *supinf*

		Player B		
		a	b	c
Player A	a	2	0	9
	b	4	4	7
	c	10	1	3

Table: What is the gain the A can be sure of obtaining ? Reminder : B wants to **minimize the gain** and acts rationally.

Pure strategy and mixed strategy

- ▶ A **pure strategy** is completely deterministic
- ▶ A **mixed strategy** assigns a probability distribution to the set of actions.

Rock Scissors Paper

- ▶ We will study mixed strategies in the Rock Scissors Paper game.

Rock Scissors Paper

- ▶ How many action tuples are possible ?

Rock Scissors Paper

Exercise 1: Action probabilities

- ▶ **cd** `minimax_and_games/zero_sum` folder.
- ▶ Modify the file **`paper_scissors_rock.py`** so that the actions performed by the two players are drawn from the relevant distributions (= strategies) **`player_X_strategy`**.
- ▶ Use can use the function **`choice`** from **`numpy`** (look for its documentation)

Rock Scissors Paper

Exercise 1 : **Action probabilities**

- ▶ Modify the file so that the correct statistics about the victory rate are computed.

Rock Scissors Paper

Exercise 2: Action probabilities

- ▶ What happens if you change the strategy of player B ?

Rock Scissors Paper

Exercise 2: Action probabilities

- ▶ What happens if you change the strategy of player B ?
- ▶ And the strategy of player A ?
- ▶ How can we interpret that result ?

Rock Scissors Paper

Let us define probabilistic **events** :

- ▶ V : A wins the game
- ▶ Br : B plays rock.
- ▶ Bs : B plays scissors.
- ▶ Bp : B plays paper.

Rock Scissors Paper

- ▶ V : A wins the game
- ▶ Br : B plays rock.
- ▶ Bs : B plays scissors.
- ▶ Bp : B plays paper.

$$\begin{aligned} P(V) &= P(V \cap (Br \cup Bs \cup Bp)) \\ &= P((V \cap Br) \cup (V \cap Bs) \cup (V \cap Bp)) \end{aligned} \tag{1}$$

Symbols :

- ▶ P means "probability of".
- ▶ \cup means "or"
- ▶ \cap means "and"

Rock Scissors Paper

$$\begin{aligned}P(V) &= P(V \cap (Br \cup Bs \cup Bp)) \\&= P((V \cap Br) \cup (V \cap Bs) \cup (V \cap Bp)) \\&= P(V \cap Br) + P(V \cap Bs) + P(V \cap Bp) \text{ (Incompatibility of events)} \\&\quad (2)\end{aligned}$$

Rock Scissors Paper : conditional probabilities

$$\begin{aligned}P(V) &= P(V \cap (Br \cup Bs \cup Bp)) \\&= P((V \cap Br) \cup (V \cap Bs) \cup (V \cap Bp)) \\&= P(V \cap Br) + P(V \cap Bs) + P(V \cap Bp) \\&= P(V|Br)P(Br) + P(V|Bs)P(Bs) + P(V|Bp)P(Bp)\end{aligned}\tag{3}$$

$P(V|Br)$ means "Probability that A wins, given that B plays rock".

Rock Scissors Paper : conditional probabilities

$P(V|Br)$ means "Probability that A wins, given that B plays rock".

If the strategy of A is $[\frac{1}{3}, \frac{1}{3}, \frac{1}{3}]$, then

- ▶ $P(V|Br) = ?$
- ▶ $P(V|Bs) = ?$
- ▶ $P(V|Bp) = ?$

Rock Scissors Paper : conditional probabilities

$P(V|Br)$ means "Probability that A wins, given that B plays rock".

If the strategy of A is $[\frac{1}{3}, \frac{1}{3}, \frac{1}{3}]$, then

- ▶ $P(V|Br) = \frac{1}{3}$
- ▶ $P(V|Bs) = \frac{1}{3}$
- ▶ $P(V|Bp) = \frac{1}{3}$

Rock Scissors Paper : conditional probabilities

$P(V|Br)$ means "Probability that A wins, given that B plays rock".

If the strategy of A is $[\frac{1}{3}, \frac{1}{3}, \frac{1}{3}]$, then

$$P(V|Br) = P(V|Bs) = P(V|Bp) = \frac{1}{3}.$$

Hence :

$$\begin{aligned} P(V) &= P(V|Br)P(Br) + P(V|Bs)P(Bs) + P(V|Bp)P(Bp) \\ &= \frac{1}{3}P(Br) + \frac{1}{3}P(Bs) + \frac{1}{3}P(Bp) \\ &= \frac{1}{3}(P(Br) + P(Bs) + P(Bp)) \end{aligned} \tag{4}$$

Rock Scissors Paper : conditional probabilities

$P(V|Br)$ means "Probability that A wins, given that B plays rock".

If the strategy of A is $[\frac{1}{3}, \frac{1}{3}, \frac{1}{3}]$, then

$$P(V|Br) = P(V|Bs) = P(V|Bp) = \frac{1}{3}.$$

Hence :

$$\begin{aligned} P(V) &= P(V|Br)P(Br) + P(V|Bs)P(Bs) + P(V|Bp)P(Bp) \\ &= \frac{1}{3}P(Br) + \frac{1}{3}P(Bs) + \frac{1}{3}P(Bp) \\ &= \frac{1}{3}(P(Br) + P(Bs) + P(Bp)) \\ &= \frac{1}{3} \end{aligned} \tag{5}$$

Biased game

Exercise 3: Alternative game with the well.

- ▶ Modify the file **paper_scissors_rock_well.py** so that the actions performed by the two players are drawn from the relevant distributions (= strategies) **player_X_strategy**, and so that the statistics are correctly computed.
- ▶ find a strategy that gives a better victory rate for A.

Statistics on strategies

Exercise 4: Learning a strategy

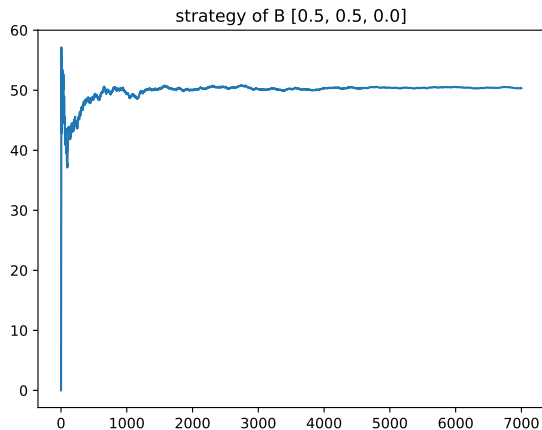
- ▶ Modify the file **paper_rock_scissors_learn.py** in order to **learn the strategy of B**, and adapt the strategy of player *A* in order to have a better victory rate.

Statistics on strategies


Exercice 5 : Learning a strategy

- ▶ Modify the file **paper_rock_scissors_learn.py** in order to **learn the strategy of B**, and adapt the strategy of player *A* in order to have a better victory rate for *A*.
- ▶ Several solutions are possible.
- ▶ You can work in groups.

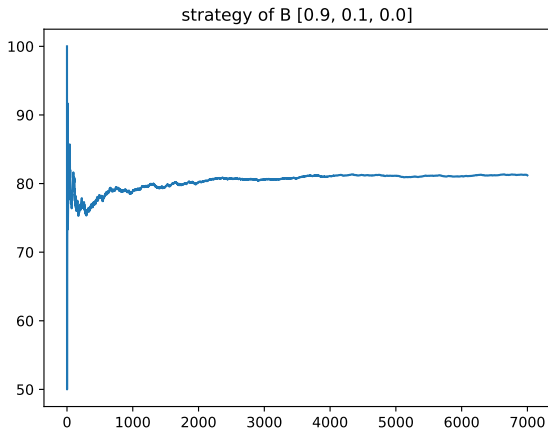
Percentage of victory



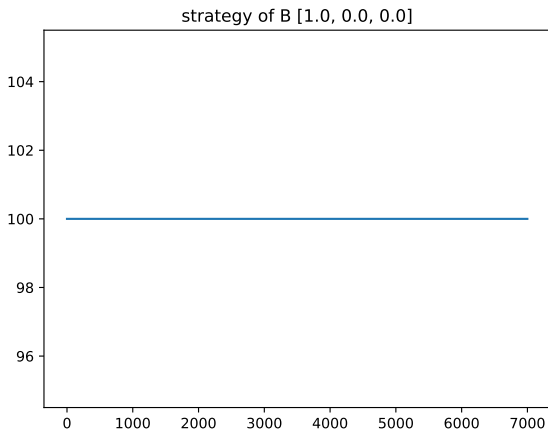
Percentage of victory



Percentage of victory



Percentage of victory



Sequential games

- ▶ We will now change the type of games studied
- ▶ We still have two players but the game consists in a **sequence of actions**, instead of a **single action**.

Sequential games

- ▶ We will now change the type of games studied
- ▶ We still have two players but the game consists in a **sequence of actions**, instead of a **single action**.
- ▶ The two players play successively, taking into account the previous actions, and also the following actions from their opponent.

Sequential games

- ▶ We will now change the type of games studied
- ▶ We still have two players but the game consists in a **sequence of actions**, instead of a **single action**.
- ▶ The two players play successively, taking into account the previous actions, and also the following actions from their opponent.
- ▶ until the game reaches a **final state**. When the game is in its final state, the players receive a score.

Sequential games

- ▶ One player is called the **maximiser**, the other player the **minimiser**.

Sequential games

- ▶ One player is called the **maximiser**, the other player the **minimiser**.
- ▶ The maximiser tries to get the **highest score**, while the minimiser tries to get the **lowest score**.

The Minimax algorithm

- ▶ We will study and implement an algorithm that computes the values of the actions of the two players.

The Minimax algorithm

- ▶ We will study and implement an algorithm that computes the values of the actions of the two players.
- ▶ **Very important hypotheses :** the agents are assumed to behave **rationally**

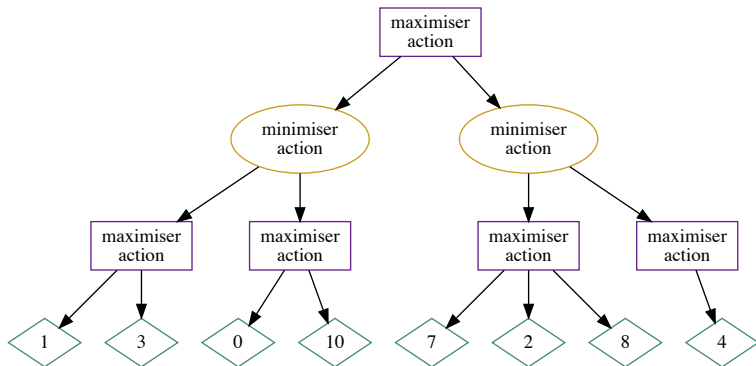


Figure: Representation of the game

Recursion

- ▶ The Minimax algorithm is based on a concept called **recursion**.

Recursion

- ▶ **Proposed definition** : a method to solve a problem based on smaller instances of the same problem.

First Recursion example

- ▶ **cd recursion**
- ▶ Please modify **factorial_rec.py** so that it computes the factorial
- ▶ $n! = 1 \times 2 \times \dots \times n$

Recursion

A recursive function always has :

- ▶ a base case
- ▶ a recursive case

Warning

- ▶ Decrease does not mean terminate !
- ▶ What happens with the example **bad_recursion** ?
- ▶ In python, you can see the recursion limit with **sys.getrecursionlimit()**

Minimax algorithm

- ▶ The **Minimax algorithm** is a recursive algorithm that computes the values of all the nodes.

Minimax algorithm

- ▶ The **Minimax algorithm** is a recursive algorithm that computes the values of all the nodes.
- ▶ It does so by propagating the information from the **leaf nodes** (the **final states of the game**) to the parent nodes.

Minimax algorithm

- ▶ The **Minimax algorithm** is a recursive algorithm that computes the values of all the nodes.
- ▶ It does so by propagating the information from the **leaf nodes** (the **final states of the game**) to the parent nodes.
- ▶ Let us apply it on an example.


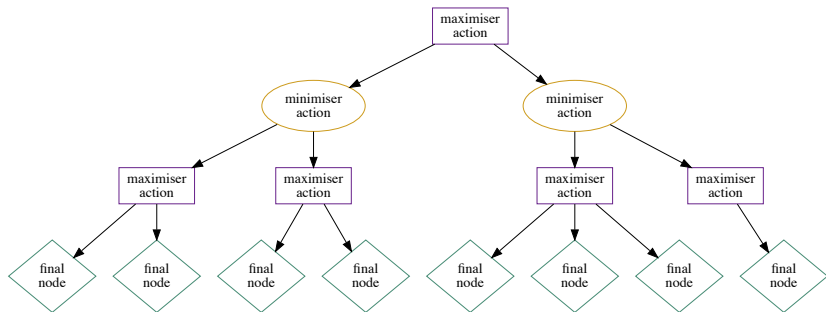
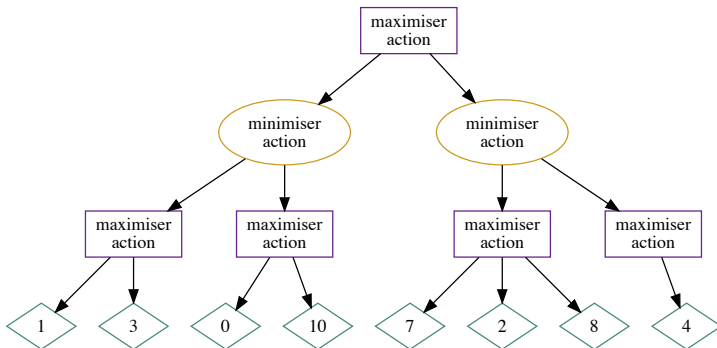


Figure: The numbers do not represent the values here : they represent the index of the node.









Figure: Values of the final states

Games and the Minimax algorithm

└ The Minimax Algorithm

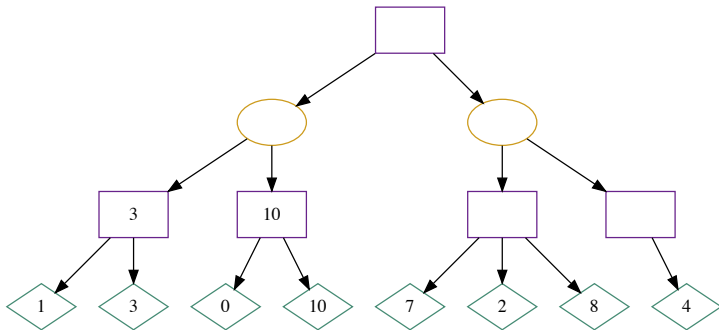
└ Minimax algorithm



Games and the Minimax algorithm

└ The Minimax Algorithm


└ Minimax algorithm



Games and the Minimax algorithm

└ The Minimax Algorithm


└ Minimax algorithm

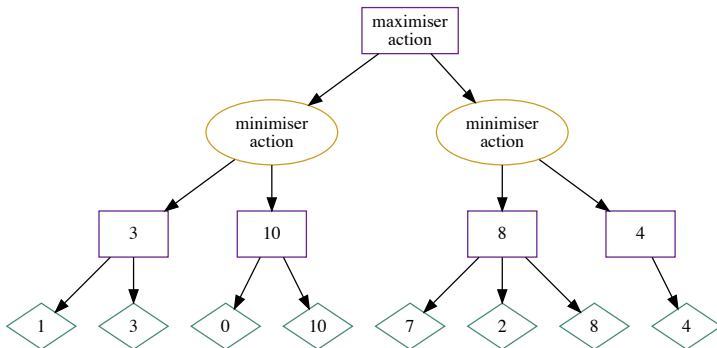


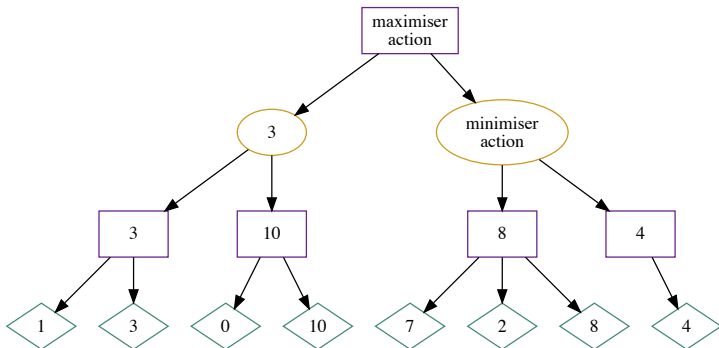
Games and the Minimax algorithm


└ The Minimax Algorithm

└ Minimax algorithm





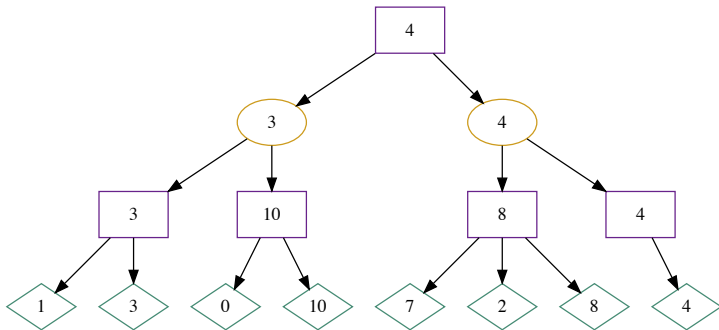




Games and the Minimax algorithm

└ The Minimax Algorithm

└ Minimax algorithm



Python dictionaries

- ▶ **dictionaries** are a useful data structure.
- ▶ Demo with **ipython**

Minimax algorithm

Exercise 6 : Implementing the algorithm

- ▶ Please use the file **minimax.py** in order to implement the algorithm.
- ▶ I inserted 2 errors in the **minimax** function.
- ▶ you can also try with different values for the final states.

Limitations of the method

- ▶ What could be the problems with the Minimax algorithm ?

Limitations of the method

- ▶ Let p be the **branching factor** of the tree. (Here, the average number of children at each node)
- ▶ Let d be the **depth** of the tree.

Limitations of the method


- ▶ Let p be the **branching factor** of the tree of actions. (Here, the average number of children at each node)
- ▶ Let d be the **depth** of the tree.
- ▶ What is the order of magnitude of the number of nodes in the tree ?

Limitations of the method

- ▶ Let p be the **branching factor** of the tree of actions. (Here, the average number of children at each node)
- ▶ Let d be the **depth** of the tree.
- ▶ What is the order of magnitude of the number of nodes in the tree ?
- ▶ So in order to run the minimax algorithm needs to perform p^d evaluations.


Exercise

- ▶ In the **Othello** game, the average number of actions in each state is around 8.
- ▶ We assume that the evaluation for one leaf node takes 1×10^{-6} seconds.




Exercise

- ▶ In the **Othello game**, the average number of actions in each state is around 8.
- ▶ We assume that the evaluation for one leaf node takes 1×10^{-6} seconds.
- ▶ How long would be the search of the minimax if we look 10 actions ahead ?



Exercise

- ▶ In the **Othello game**, the average number of actions in each state is around 8.
- ▶ We assume that the evaluation for one leaf node takes 1×10^{-6} seconds.
- ▶ This duration is too long to be used.



Conclusion

For real games, we need to use either :

- ▶ a smaller tree (look not too many steps ahead)
- ▶ a faster algorithm

Alpha-Beta pruning

- ▶ We will study a method that optimizes the Minimax algorithm.

Alpha-Beta pruning

- ▶ We will study a method that optimizes the Minimax algorithm.
- ▶ It does so by preventing us from computing **useless nodes**

Alpha-Beta pruning

- ▶ We will study a method that optimizes the Minimax algorithm.
- ▶ It does so by preventing us from computing **useless nodes**
- ▶ Let us do it on an example.

Alpha-Beta pruning






Figure: We use different final state values


Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning




Alpha-Beta pruning



Alpha-Beta pruning



Alpha-Beta pruning



Implementation

Exercise 7: Alpha beta pruning

- ▶ Let us now implement the algorithm
- ▶ use the file **alpha_beta.py** in order to implement the algorithm.
- ▶ There are several mistakes in the code.

Verification

- ▶ Please verify that the AlphaBeta algorithm gives the same result as the Minimax algorithm !

Other values

- ▶ Please try to modify the initial values to change the behavior of the algorithm.

Orders of magnitude

- ▶ If N is the number of nodes explored by the normal minimax algorithm, the number of nodes explored by Alphabeta can be of order of magnitude \sqrt{N}
- ▶ This is a great improvement.
- ▶ **However**, please note that the improvement depends on the tree. Sometimes the pruning will not accelerate the algorithm that much.

Intermediate scoring

- Sometimes it is not possible to explore the entire tree, if it is too large, even with alpha beta pruning.

Intermediate scoring

- ▶ Sometimes it is not possible to explore the entire tree, if it is too large, even with alpha beta pruning.
- ▶ In this situation, it is necessary to use **intermediate scoring**

Intermediate scoring

- ▶ Sometimes it is not possible to explore the entire tree, if it is too large, even with alpha beta pruning.
- ▶ In this situation, it is possible to use **intermediate scoring**
- ▶ It is a heuristic : there is no theoretical proof that it yields the best solution, but it permits computation

Intermediate scoring and phantom of the opera

- Can you think of an intermediate scoring ?



Intermediate scoring and phantom of the opera

- ▶ Can you think of an intermediate scoring ?
- ▶ What are the depth and the width of the tree ?

